[3 nodes Kafka cluster]

1. Install kafkacat
    a. `brew install kcat`
2. Create a `docker-compose.yml` file
3. This should ensure that Zookeeper always starts before the Kafka server and ends after it
4. `docker-compose.yml`
    a. Should have 6 services - 3 zookeeper services and 3 kafka services
    b. **`image: confluentinc/cp-kafka:latest`** is the docker image for deploying and running of the community version of Kafka
    c. **`image: confluentinc/cp-zookeeper:latest`** is the docker image for deploying and running Zookeeper
    d. Ensure that the service names and KAFKA_BROKER_ID are unique
    e. **`depends_on:`** Kafka depends on Zookeeper to run, so the Zookeeper keys are included in this property to ensure Docker will start the Zookeepers first before the Kafka servers
    f. Zookeeper 1, 2 and 3 are listening to connections by client such as Kafka servers on port 2181 by the ZOOKEEPER_CLIENT_PORT property
    g. All servers must expose unique ports to the host machine
        i. The 3 zookeeper servers are exposed to the host at port 22181, 32181 and 42181 respectively by the ports property
        ii. Same for kafka servers where they will be listening on ports 29092, 39092 and 49092 respectively
    h. Every machine in the Zookeeper ensemble should know about other machines in the ensemble using the ZOOKEEPER_SERVERS property
        i. **`ZOOKEEPER_SERVERS: zookeeper1:2888:3888; zookeeper2:2888:3888;zookeeper3:2888:3888;`**

docker-compose.yml > { } services > { } kafka1 > [ ] ports > 0

docker-compose.yml (compose-spec.json)

```yaml
version: '2'
services:
  zookeeper1:
    image: confluentinc/cp-zookeeper:latest
    hostname: zookeeper1
    ports:
      - "22181:22181"
    environment:
      ZOOKEEPER_SERVER_ID: 1
      ZOOKEEPER_CLIENT_PORT: 22181
      ZOOKEEPER_TICK_TIME: 2000
      ZOOKEEPER_INIT_LIMIT: 5
      ZOOKEEPER_SYNC_LIMIT: 2
      ZOOKEEPER_SERVERS: zookeeper1:2888:3888;zookeeper2:2888:3888;zookeeper3:2888:3888
    networks:
      - proxy

  zookeeper2:
    image: confluentinc/cp-zookeeper:latest
    hostname: zookeeper2
    ports:
      - "32181:32181"
    environment:
      ZOOKEEPER_SERVER_ID: 2
      ZOOKEEPER_CLIENT_PORT: 32181
      ZOOKEEPER_TICK_TIME: 2000
      ZOOKEEPER_INIT_LIMIT: 5
      ZOOKEEPER_SYNC_LIMIT: 2
      ZOOKEEPER_SERVERS: zookeeper1:2888:3888;zookeeper2:2888:3888;zookeeper3:2888:3888
    networks:
      - proxy

  zookeeper3:
    image: confluentinc/cp-zookeeper:latest
    hostname: zookeeper3
    ports:
      - "42181:42181"
    environment:
      ZOOKEEPER_SERVER_ID: 3
      ZOOKEEPER_CLIENT_PORT: 42181
      ZOOKEEPER_TICK_TIME: 2000
      ZOOKEEPER_INIT_LIMIT: 5
      ZOOKEEPER_SYNC_LIMIT: 2
      ZOOKEEPER_SERVERS: zookeeper1:2888:3888;zookeeper2:2888:3888;zookeeper3:2888:3888
    networks:
      - proxy
```

```
  kafka1:
    image: confluentinc/cp-kafka:latest
    hostname: kafka1
    ports:
      - "29092:29092"
    depends_on:
      - zookeeper1
      - zookeeper2
      - zookeeper3
    environment:
      KAFKA_BROKER_ID: 1
      KAFKA_ZOOKEEPER_CONNECT: zookeeper1:22181,zookeeper2:22181,zookeeper3:22181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka1:29092
    networks:
      - proxy

  kafka2:
    image: confluentinc/cp-kafka:latest
    hostname: kafka2
    ports:
      - "39092:39092"
    depends_on:
      - zookeeper1
      - zookeeper2
      - zookeeper3
    environment:
      KAFKA_BROKER_ID: 2
      KAFKA_ZOOKEEPER_CONNECT: zookeeper1:22181,zookeeper2:22181,zookeeper3:22181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka2:39092
    networks:
      - proxy

  kafka3:
    image: confluentinc/cp-kafka:latest
    hostname: kafka3
    ports:
      - "49092:49092"
    depends_on:
      - zookeeper1
      - zookeeper2
      - zookeeper3
    environment:
      KAFKA_BROKER_ID: 3
      KAFKA_ZOOKEEPER_CONNECT: zookeeper1:22181,zookeeper2:22181,zookeeper3:22181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://kafka3:49092
    networks:
      - proxy
networks:
  proxy:
    driver: bridge
```

5. Update the /etc/hosts file
    a. `sudo nano /etc/hosts`
    b. Enter password
    c. Add the following line
        i. `0.0.0.0 kafka1 kafka2 kafka3`

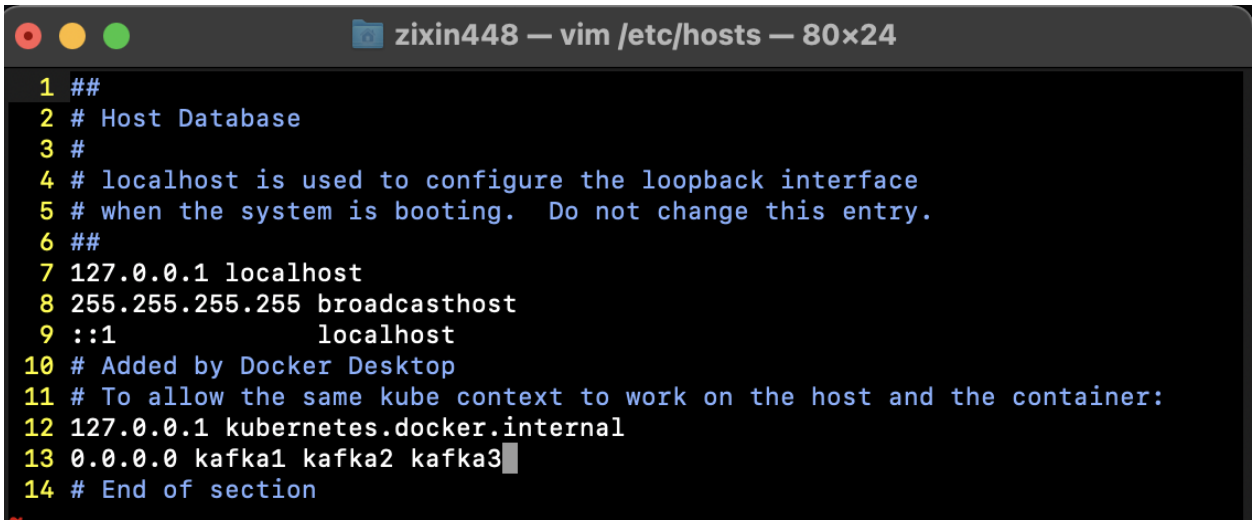

6. Start the cluster using the docker compose command
    a. `docker-compose up -d`

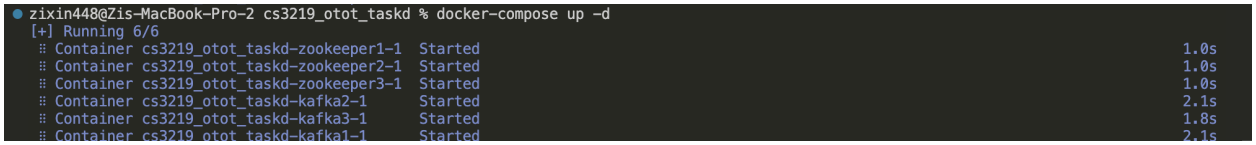

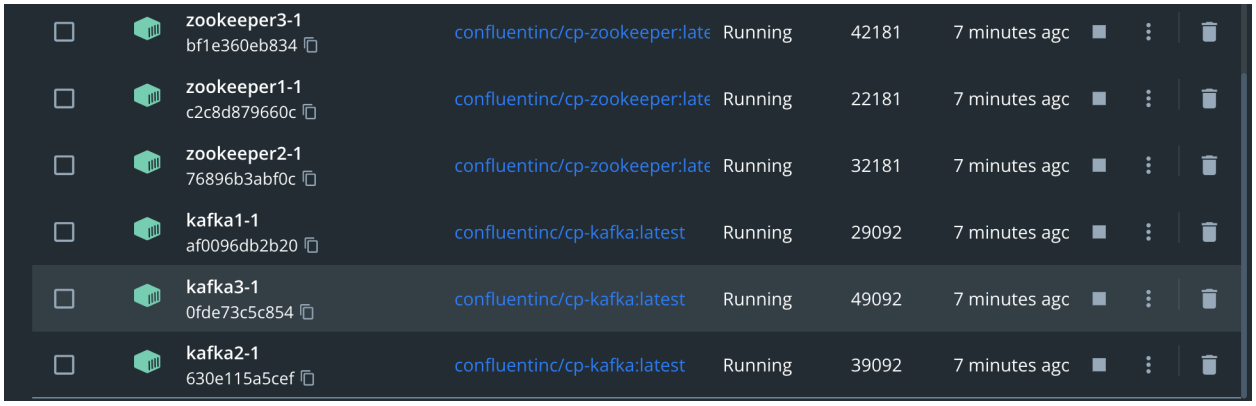

7. Check that the services are running
    a. `docker-compose ps`
8. Pick a controller
    a. `sudo kcat -L -b kafka1:29092` (or kafka:39092 or kafka:49092)
    b. Enter password

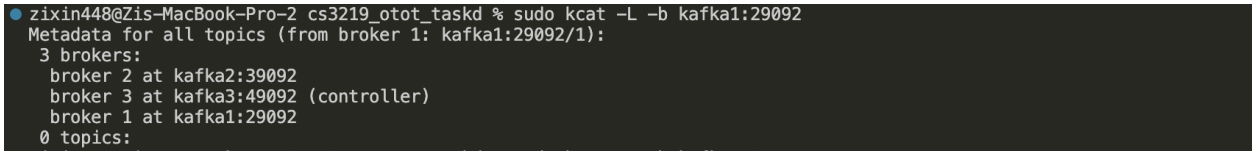

9. Create a topic
    a. `docker run --net=host --rm confluentinc/cp-kafka:latest kafka-topics --create --topic mytopic --partitions 1 --replication-factor 3 --if-not-exists --bootstrap-server localhost:29092`

```
zixin448@Zis-MacBook-Pro-2 cs3219_otot_taskd % docker run --net=host --rm confluentinc/cp-kafka:latest kaf
ka-topics --create --topic mytopic --partitions 1 --replication-factor 3 --if-not-exists --bootstrap-serve
r localhost:29092
```

10. Open new terminal and run a kafka server as producer
    a. `kcat -P -b kafka1:29092 -t mytopic`
    b. -P: flag for producer
    c. -t: topic name flag
    d. -b: broker chosen, in this case is kafka1 server
11. In a separate terminal, run another server as consumer
    a. `kcat -C -b kafka2:39092 -t mytopic`
    b. -C: flag for consumer
12. Check the Pub-Sub messaging between the two terminals
    a. Send messages at producer terminal, click `return` to enter next line
    b. When done, `Ctrl-D` to publish
    c. Visit consumer terminal to see messages published by producer

```
zixin448@Zis-MacBook-Pro-2 cs3219_otot_taskd % kcat -P -b kafka1:29092 -t mytopic
hello do you receive
testing 2
zixin448@Zis-MacBook-Pro-2 cs3219_otot_taskd %
```

```
zixin448@Zis-MacBook-Pro-2 cs3219_otot_taskd % kcat -C -b kafka2:39092 -t mytopic
% Reached end of topic mytopic [0] at offset 0
hello do you receive
testing 2
% Reached end of topic mytopic [0] at offset 2
```